

Grinta for Life – Founder Agent

Technical Specification – v1.0

Five-layer information architecture for Jay's drafting agent

Author: Jay Crain

Date: May 2026

Version: 1.0

Status: Active. v1.0 applies the May 2026 language cascade (Member, asset, ID Score, Pulse Check, the four IDQ dimensions Physical/Self/Social/Outlook, IDQ intake & retake terminology, single October platform launch model, bands retired). Working specification for developer hire, agency RFP, and internal team alignment.

Companion documents: Founder Agent Operating Moments SVG (front-page diagram); Founder Agent Data Model SVG; Member Agent Technical Specification v1.1; AI Governance Framework v1.2; Brand Standards v1.3; Tom Miller Member Journey v4.3; Reshma Patel Member Journey v1.1; Martin Jackson Member Journey v1.1; Platform Development Plan v4.

How to read this Spec

This document is the working technical specification for the Founder Agent – the drafting AI that produces Jay's Member-facing correspondence and surfaces Jay's voice across the dashboard's educational layer. It describes the five layers of information the agent operates on, the six operating moments where the agent produces output, the voice and governance enforcement that constrains the agent, the Jay-review-before-send mandate that gates every Member-facing email, the failure modes the architecture mitigates, and the build sequence to platform launch in mid-October 2026.

The audience is multi-purpose. A developer being hired reads this as the hand-off. An agency bidding on the build reads it as the basis for a statement of work. The internal team reads it as the architectural reference. The same content serves all three.

The Founder Agent is architecturally simpler than the Member Agent in surface count and conversation pattern, but governance-bottlenecked at every event-driven output. Where the Member Agent has a continuous-companion architecture with member-led conversation, the Founder Agent is event-driven with Jay as the human in the loop on every send. The Spec is built to make that governance posture legible to every reader.

The Spec is built for iteration. The product experience team is iterating the three canonical Member journeys (Tom Miller v4.3, Reshma Patel v1.1, Martin Jackson v1.1); when the journeys update, sections of this Spec update with them. Section 9 names every journey dependency explicitly. Versioning convention: Journey v1.x cascade maps to Founder Spec v1.x.

G4L Founder Agent — Operating Moments

When Jay's voice surfaces — companion to the Founder Agent Data Model

Five event-driven moments plus one ambient educational-layer surfacing — every event-driven draft is Jay-review-before-send

1. Post-IDQ Welcome

Once per cycle — after IDQ intake submission

GOVERNANCE
JAY-REVIEW: YES

TRIGGER

Member submits IDQ intake (Cycle 1) or esms/different IDQ intake (Cycle 2 onset)

INPUTS

IDQ intake responses - named Door - intake quote
Selected track length - noun (if proposed) - Jay's voice corpus

POSTURE

Acknowledge the quote in plain language - name the Door
Frame the cycle ahead - no metrics or numbers

OUTPUTS

Drafted welcome email - subject + body + sign-off
Tonal calibration notes - Member-specific phrasing

HAND-OFFS

⌘ Jay reviews + approves + sends - ⌘ Member email inbox
⌘ Member Agent context (welcome sent timestamp)

2. Retake Commentary

After every IDQ retake — 30 day cadence

GOVERNANCE
JAY-REVIEW: YES

TRIGGER

Member completes IDQ retake
Member Agent has surfaced retake reflection

INPUTS

Retake numbers - signed deltas - direction arrows
Recent Pulse Check texture - asset closing lines - Jay's corpus

POSTURE

Name what the chart shows in Member-specific terms
Reference Door - quote Member's own words back - no false praise

OUTPUTS

Drafted retake email - subject + body
Optional milestone callout if Phase boundary crossed

HAND-OFFS

⌘ Jay reviews + approves + sends - ⌘ Member email inbox
⌘ Member Agent context (retake email sent)

3. Milestone Commentary

At specific moments — Phase transitions, asset closures, behavior milestones

GOVERNANCE
JAY-REVIEW: YES

TRIGGER

Gating asset closes that marks a Phase boundary OR
Member-named milestone surfaces (e.g., body or behavior change)

INPUTS

Milestone signal - Member context - Door - Jay's corpus
Recent Conversation Surface excerpts - Pulse Check signals

POSTURE

Acknowledge specifically - not generically
Tie to Door - tie to noun if named - measured tone

OUTPUTS

Drafted milestone email - subject + body
Optional Circle prompt if Member opts in

HAND-OFFS

⌘ Jay reviews + approves + sends - ⌘ Member email inbox
⌘ Member Agent context (milestone email sent)

4. False Start Return

Event-driven — when Member returns from missed cadence window

GOVERNANCE
JAY-REVIEW: YES

TRIGGER

Member missed Pulse Check or IDQ window AND
returns within recovery window (Member Agent flag)

INPUTS

Last activity timestamp - miss reason if known
Member's last reflection - Jay's corpus - Door

POSTURE

No re-engagement framing - no recovery narrative
Steady welcome back - reference what was happening before

OUTPUTS

Drafted return email - short, light touch
Optional one-line SMS draft for opt-in Members

HAND-OFFS

⌘ Jay reviews + approves + sends - ⌘ Member email inbox
⌘ Member Agent context (cadence resumed)

5. Cycle 2 Welcome

Once per Cycle 2 onset — earned or offered

GOVERNANCE
JAY-REVIEW: YES

TRIGGER

Cycle 2 earned via trigger event OR
Cycle 2 offered after stable Cycle 1 close

INPUTS

Cycle 2 IDQ intake - new Door (if renamed)
Noun (compound if added) - full Cycle 1 arc - Jay's corpus

POSTURE

Name what changed since Cycle 1 close - name new Door
Frame Cycle 2 length and pacing - honor the choice to return

OUTPUTS

Drafted Cycle 2 welcome email - subject + body
Track length recommendation if Member is open to it

HAND-OFFS

⌘ Jay reviews + approves + sends - ⌘ Member email inbox
⌘ Member Agent context (Cycle 2 active)

6. Ambient Educational Surfacing

Always available — surfaces on dashboard drill-down and asset context

GOVERNANCE
JAY-REVIEW: N/A (library)

TRIGGER

Member drills into a dimension on dashboard OR
Member opens an asset that surfaces Founder framing

INPUTS

Door - starting profile - current dimension state
Pre-approved Founder-voice library - Brand Standards v1.3

POSTURE

No per-Member drafting - no live generation
Surface library content shaped to Door + dimension

OUTPUTS

Door-specific framing card - dimension framing card
Pulled from library, not generated

HAND-OFFS

⌘ Dashboard educational layer renders the card
⌘ Asset view embeds framing in Founder voice register

Governance principle — Jay-review-before-send

The five event-driven operating moments produce drafts. The drafts do not send. Jay reads, approves, and sends every Member-facing email personally.

Draft ⌘ review ⌘ send

Every event-driven moment produces a draft.
Jay reads, approves, sends. No exceptions.

Member Agent context updated

After Jay sends, the Member Agent records the send and surfaces it in next reflection.

Ambient layer is library, not draft

Educational-layer cards are pre-approved.
No live generation. No per-Member drafting.

1. Executive Architecture Summary

1.1 The Founder Agent's role

The Founder Agent is Jay Crain's drafting agent. Its job is to produce Member-facing correspondence – emails, milestone notes, retake commentaries, Cycle 2 welcomes – in Jay's voice, shaped to each individual Member's context, and to surface that voice in the dashboard's educational layer when Members drill into a dimension or open an asset that wants Founder framing.

The Founder Agent does not coach, prescribe, send, or substitute for Jay's personal judgment. It drafts. Jay reads, approves, and sends every Member-facing email personally. The Founder Agent does not have a send button. This is not a technical limitation; it is the canonical governance posture that defines the agent.

The agent's most consequential function is being a faithful Jay-voice drafter at scale. Without the agent, Jay personally writes every welcome, every retake commentary, every milestone note, every Cycle 2 welcome for every Member – a workload that does not survive past Member count fifty. With the agent, Jay reviews drafts, edits where needed, and sends. The architecture is built to make Jay's voice durable across thousands of Members without diluting it or substituting for it.

1.2 Reference architecture

The Founder Agent runs on the same platform reference architecture as the Member Agent, defined in Appendix C v1.3:

- Web app framework + hosting: Next.js on Vercel
- Data layer: Supabase (bundled Postgres, Auth, Storage, Realtime)
- Authentication: Supabase Auth (bundled with data layer)
- AI provider: Anthropic API (Claude family); provider abstraction layer in place for hedge
- CRM spine: HubSpot (relationship view; every Member, every campaign, every email)
- Email provider: HubSpot's transactional email rail (existing infrastructure)
- Review interface: Founder Agent Review UI – a Jay-only surface where drafts are read, edited, approved, and sent

1.3 The five-layer information model

The Founder Agent operates on a five-layer information model. Layer 1 is the agent's identity and configuration – system prompt, Jay-voice fingerprint, governance enforcement rules. Layer 2 is Jay's writing corpus – prior emails, second-edition Part Three material, BP front matter, every artifact Jay has personally written or signed off on that defines the Jay voice. Layer 3 is the Member Profile mirror – a read-only projection of the Member Agent's Layer 3 (Member Profile) data, scoped to fields the Founder Agent needs for drafting. Layer 4 is the Member History mirror – a read-only projection of the Member Agent's Layer 4 (Member History) data, scoped to recent retakes, recent Pulse Check signals, recent asset closures, recent Member-Agent Conversation

Surface excerpts. Layer 5 is the Draft Output – the per-Member, per-moment artifact the agent produces and writes to the `founder_agent_drafts` table for Jay's review.

The five-layer model differs from the Member Agent's six-layer model in three ways. First, the Founder Agent has no Layer 6 – it does not have a Conversation Surface and does not produce Member-facing conversation. The drafted output is the agent's only output. Second, Layers 3 and 4 are read-only mirrors rather than primary stores – the Member Agent owns the canonical Member data, and the Founder Agent reads it through a scoped view. Third, Layer 2 (Jay's writing corpus) is materially different in character from the Member Agent's Layer 2 (Program Knowledge Corpus) – the Founder Agent's Layer 2 is a voice fingerprint and corpus, not a knowledge base.

1.4 The six operating moments

The Founder Agent operates at six distinct moments. Five are event-driven – they produce a draft for Jay-review-before-send. One is ambient – it surfaces pre-approved library content without per-Member drafting.

1. *Post-IDQ Welcome* – once per cycle, after IDQ intake submission. Produces a welcome email draft.
2. *Retake Commentary* – after every IDQ retake (60-day cadence). Produces a retake email draft.
3. *Milestone Commentary* – at gating asset closures that mark Phase boundaries or at Member-named milestones. Produces a milestone email draft.
4. *False Start Return* – event-driven, when a Member returns from a missed cadence window. Produces a return email draft.
5. *Cycle 2 Welcome* – once per Cycle 2 onset (earned or offered). Produces a Cycle 2 welcome email draft.
6. *Ambient Educational Surfacing* – always available. Surfaces pre-approved Founder-voice library content on the dashboard educational layer when Members drill into a dimension or open an asset with Founder framing.

The five event-driven moments are governance-bottlenecked: every draft is Jay-review-before-send. The ambient layer is library-based: pre-approved content surfaces without per-Member drafting and therefore without per-event review. Section 4 details each moment.

1.5 Versioning convention

This Spec is versioned in lockstep with the canonical Member journeys. The current canonical journey set is Tom Miller v4.3, Reshma Patel v1.1, and Martin Jackson v1.1. When those journeys cascade to a new version (next cascade or new persona addition), this Spec's version bumps. v1.0 maps to the current journey set.

Decision provenance: every architectural decision in this Spec carries a comment naming the decision source. Sources are one of: the May 2026 language cascade brainstorm, the AI Governance Framework v1.2, the Brand Standards v1.3 Inputs, the Member Agent Tech Spec v1.1, the Tom v4.3 master narrative, or Jay's direct authorship.

2. Reference Architecture

2.1 Position relative to the Member Agent

The Founder Agent and the Member Agent share infrastructure but serve different functions. The Member Agent is continuous, conversational, and Member-led – it talks with the Member, holds the Member's longitudinal record, and surfaces the framework's reading of the Member's trajectory. The Founder Agent is event-driven, draft-producing, and Jay-gated – it produces Member-facing correspondence on Jay's behalf and surfaces Jay's voice in the dashboard's educational layer.

The two agents communicate through shared data. The Member Agent writes Layer 3 (Member Profile) and Layer 4 (Member History) records to Supabase. The Founder Agent reads scoped views of those records. When the Founder Agent produces a draft, it writes the draft to the founder_agent_drafts table. When Jay reviews, edits, and sends, the send event is captured in the founder_agent_drafts table and surfaced back to the Member Agent's context so that the Member Agent's next reflection can reference what Jay just sent.

Founder Agent and Member Agent – role comparison

Dimension	Member Agent	Founder Agent
Pattern	Continuous companion. Member-led conversation. Always available.	Event-driven. Agent-led drafting. No conversation surface.
Surface to Member	Conversation Surface (chat UI). Dashboard reflections. Asset reflections.	Email (drafted, Jay-sent). Dashboard educational layer (library).
Output type	Live conversational responses; reflective synthesis; behavioral signal capture.	Drafted emails; library-sourced dashboard cards.

Governance	AI disclosure as first line. Crisis disclosure routing. Greg sub-corpus boundaries.	Jay-review-before-send mandatory on every event-driven draft. Library-only on ambient layer.
Data ownership	Owens Layer 3, Layer 4, Layer 6.	Reads scoped mirrors of Member Agent Layer 3 and Layer 4. Owns Layer 2 (Jay's voice corpus) and Layer 5 (Drafts).
Frequency	Every Member session, every retake, every Pulse Check, every asset.	Per-Member: 4-6 event-driven sends per cycle. Per-Member-drilldown: ambient layer surfaces on demand.

2.2 The two data spines

Same as Member Agent Tech Spec v1.1 §2.2. The Founder Agent reads from the program spine (Supabase warehouse) and writes its outputs to the program spine. The HubSpot relationship spine is the rail through which Jay-approved drafts become actual emails.

HubSpot – relationship spine

Every Member is a HubSpot contact. Every Member-facing email Jay sends – whether drafted by the Founder Agent or written by Jay personally – is sent via HubSpot's transactional email rail and logged on the contact record. The Founder Agent does not write to HubSpot directly. Jay-approved sends move from the founder_agent_drafts table to HubSpot via the Review UI's send action.

Supabase warehouse – program spine

Every Founder Agent draft is a row in the founder_agent_drafts table. Every draft carries the Member ID, the operating moment, the input snapshot, the generated draft text, Jay's edits if any, the approval timestamp, and the send timestamp. The founder_agent_library table holds the pre-approved Founder-voice content that the ambient layer surfaces.

2.3 Provider selection – Anthropic

Same provider selection rationale as the Member Agent Tech Spec v1.1 §2.3: Anthropic's Claude family chosen on voice quality, instruction-following fidelity for the Jay-review-before-send governance posture, and the model's demonstrated capacity for long-context drafting with corpus citation. Provider abstraction layer in place for hedge. The Founder Agent's drafting pattern – short, single-response generation from a substantial corpus-grounded prompt – is well within the model's strengths.

2.4 Cost ceiling and reliability commitments

Founder Agent inference volume is materially lower than Member Agent volume. Estimated per-Member inference per cycle: 4-6 event-driven drafts in Cycle 1 (welcome, 1-6 retakes depending on track, 0-2 milestones, 0-2 returns, 1 Cycle 2 welcome). At Member count one thousand in Year 1, that is roughly four to six thousand drafts annually. Each draft is a long-context inference with grounded retrieval; cost per draft is meaningful but the total volume is bounded.

Reliability commitment: every event-driven moment that fires must produce a draft that lands in the Review UI within five minutes of the trigger. Jay's review SLA is target-of-24-hours from draft availability; the Member-facing email lands when Jay sends. The system is not on a cron; it is on a Jay-time clock.

3. Layer-by-Layer Technical Specification

3.1 Layer 1 – Founder Agent Identity & Configuration

Purpose

Defines who the Founder Agent is, what it does, what it does not do, and how it must behave. Loaded once at agent initialization. Constant across all operating moments and all Members.

Contents

- System prompt – identity, role, governance posture, Jay-review-before-send mandate as immutable rule
- Jay-voice fingerprint – the canonical Jay voice attributes: tone, register, sentence rhythm, signature lines, prohibited constructions
- Governance enforcement rules – do-not-send instruction, do-not-fabricate-Member-data instruction, do-not-substitute-for-clinical-judgment instruction
- Brand Standards v1.3 voice rules – the voice rules section of Brand Standards loaded as agent constraint
- AI Governance Framework v1.2 prohibitions – the prohibitions section loaded as agent constraint

Storage

Layer 1 lives in version control as system_prompt.md in the founder_agent repository. The system prompt is the single source of truth for the agent's identity. Updates to Layer 1 require a pull request reviewed by Jay before merging. Layer 1 changes are versioned with the Founder Spec version.

Loading

Loaded at agent initialization for every inference call. Sent as the system prompt to the Anthropic API on every request. Approximately 8,000-12,000 tokens depending on Brand Standards v1.3 expansion.

Voice and governance enforcement at this layer

Layer 1 is where the Jay-review-before-send governance is structurally enforced. The system prompt names this as an immutable rule. The agent does not have a send tool. The agent does not have a HubSpot-write tool. The architecture makes autonomous sending physically impossible – the only tool the agent has is write-to-founder_agent_drafts table. Sending requires the Review UI, which requires Jay's authentication, which is outside the agent's reach.

Failure modes

Layer 1 drift – the system prompt updated without versioning, breaking voice consistency.

Mitigation: pull-request gating on system_prompt.md, Jay review required, version tagged in the founder_agent_drafts table for every draft so any voice regression can be traced to a specific prompt version.

Dependencies

- Brand Standards v1.3 – voice rules section
- AI Governance Framework v1.2 – prohibitions section
- Jay's writing corpus (Layer 2) – referenced from Layer 1 as the voice authority

3.2 Layer 2 – Jay's Writing Corpus

Purpose

The corpus that defines Jay's voice. The agent's grounded voice authority. When the agent drafts a Member-facing email, it grounds the voice in Layer 2 retrieval rather than generating Jay-shaped prose from Layer 1 instructions alone. Layer 2 is what makes the drafts sound like Jay rather than like an Anthropic model writing in a generic motivational register.

Contents

- Prior Member-facing emails – every email Jay has personally written or sent, captured to the corpus with metadata (Member ID, operating moment, send date, redacted Member-specific identifiers)
- Second edition Part Three material – the third part of the GRINTA! second edition, which contains the largest concentration of Jay's reflective prose
- BP front matter – the founder's-voice sections of the Business Plan v6.1 (introduction, why-this-now sections, framework reflections)
- Brand Standards v1.3 voice examples – the in-line voice examples in Brand Standards v1.3 Inputs
- Jay's signed-off Circle posts – if/when Jay posts in the Circle community spaces, those posts join the corpus
- Journey master excerpts – the Founder Agent voice register sections from Tom v4.3, Reshma v1.1, and Martin v1.1 (educational-layer card prose)

Storage

Layer 2 lives in a Supabase founder_agent_corpus table, with one row per artifact. Each row carries the artifact text, the source (email/book/BP/Circle/journey), the date, a vector embedding for retrieval, and Jay's signoff status. Only signoff-status-confirmed rows are eligible for retrieval. Vector embeddings are produced at corpus ingestion using the same embedding model the Member Agent uses for its Greg sub-corpora.

Approximate corpus size at launch: 200-400 artifacts, totaling 80,000-150,000 tokens of raw text. Retrieval surfaces the top 5-10 most relevant artifacts per draft request, totaling 8,000-15,000 tokens of grounded context.

Loading and retrieval

Retrieval is per-draft. When the agent receives a draft request for an operating moment + Member, it constructs a retrieval query from the Member context (Door, dimension state, recent activity, intake quote if relevant) and the operating moment template. The retrieval pulls the most-relevant corpus rows. Retrieved artifacts are included in the inference context as labeled voice references.

Update and persistence rules

Layer 2 is append-only and Jay-curated. New artifacts join the corpus only after Jay's signoff. The corpus does not auto-learn from the Founder Agent's own drafts — doing so would create a recursive voice drift. Layer 2 grows when Jay personally adds an artifact, not when the agent ships a draft Jay approved.

Voice and governance enforcement at this layer

Layer 2 is the voice fingerprint's substrate. The system prompt's Jay-voice instructions are general; Layer 2 retrievals are specific. The agent's drafts read like Jay because the drafts are continuously grounded in artifacts Jay has personally signed off on.

Failure modes

Voice drift via corpus contamination — a non-Jay artifact admitted to the corpus shifts the retrieval pattern and degrades draft fidelity. Mitigation: signoff-status-confirmed gate on retrieval; periodic corpus audit by Jay (target quarterly); voice-quality sampling of drafts.

Retrieval miss — the corpus does not contain Jay-voice precedent for a specific Member context (e.g., a new Door not yet represented). Mitigation: corpus expansion is part of the v1.1 roadmap as new personas join the canonical set; Brand Standards v1.3 voice rules in Layer 1 provide fallback structure when retrieval is sparse.

Dependencies

- Brand Standards v1.3 — voice examples ingested
- Second edition Part Three — ingested as Jay-curated artifacts
- BP v6.1 front matter — ingested as Jay-curated artifacts
- Three canonical journeys — Founder Agent voice register sections ingested as artifacts

3.3 Layer 3 – Member Profile Mirror

Purpose

A read-only projection of the Member Agent's Layer 3 (Member Profile) data, scoped to the fields the Founder Agent needs for drafting. The Founder Agent does not write to the Member Profile; the Member Agent owns it. The Founder Agent reads through this scoped view.

Contents and schema

The Member Profile Mirror contains:

- Member identity – Member ID, first name, opt-in flags (SMS, Circle, email)
- Door – named Door for current cycle (e.g., The Career, The Aging Parents, The Body, The Outlook)
- Noun – currently-named identity noun (e.g., RUNNER, ATHLETE, ANCHOR, ATHLETE-PARTNER)
- Track length – selected at intake (3-month, 6-month, 9-month, 12-month)
- Intake quote – the Member's intake quote, verbatim
- Current cycle – Cycle 1 or Cycle 2 (or later)
- Current Phase – Foundation / Rebuild / Sustain
- Cycle 2 trigger type – earned (crisis-driven) or offered (stable Cycle 1 close); null if not yet at Cycle 2

Loading

Loaded per draft request. Read from a Supabase view (`member_profile_for_founder_agent`) that exposes only the fields above. Other Member Agent Layer 3 fields (e.g., behavioral signals, detailed onboarding chapter text) are not visible to the Founder Agent. The view's scope is defined in code and audit-loggable.

Update and persistence rules

Read-only. The Founder Agent does not write to Layer 3. The Member Agent's writes propagate via the view immediately. If the Member updates their Door or noun mid-cycle, the Founder Agent's next draft picks up the change automatically.

Voice and governance enforcement at this layer

Layer 3 is the Member-specific grounding for every draft. The drafts reference the Door, the noun, the track length, the intake quote, and the Cycle context not because the system prompt asks the agent to mention these things, but because Layer 3 is loaded as context on every draft request. The agent cannot draft a generic welcome – the Member's Door is in the context, the Member's intake quote is in the context, the track length is in the context. Generic drafts are architecturally hard to produce.

Failure modes

Stale Layer 3 – the Member Agent has updated the Door or noun but the Founder Agent's draft request runs against a cached read. Mitigation: no caching; every draft request reads fresh.

Missing Layer 3 – the Member has not yet completed IDQ intake but a draft request fires anyway. Mitigation: operating-moment triggers require IDQ intake completion as a precondition. The agent will refuse to draft without Layer 3 populated.

Dependencies

- Member Agent Tech Spec v1.1 §3.3 – the source of Layer 3 in the Member Agent
- Supabase view definition – `member_profile_for_founder_agent`

3.4 Layer 4 – Member History Mirror

Purpose

A read-only projection of the Member Agent's Layer 4 (Member History) data, scoped to the fields the Founder Agent needs for drafting. The Founder Agent needs recent retake numbers, recent Pulse Check signals, recent asset closures, and recent Member-Agent Conversation Surface excerpts. The Founder Agent does not need every behavioral signal the Member Agent tracks.

Contents and schema

The Member History Mirror contains, per Member, the last 90 days of:

- *IDQ retake records* – retake number, date, raw ID Score, normalized ID Score, the four dimension scores (Physical, Self, Social, Outlook), signed deltas from intake and previous retake
- *Pulse Check signals* – Pulse Check date, dimension focus, signal magnitude, brief texture if the Member Agent generated a texture note
- *Asset closures* – asset name, closure date, closing reflection field text (the Member's own words at asset close)
- *Conversation Surface excerpts* – quoted excerpts the Member Agent flagged as relevant for Founder Agent draft context (e.g., the Tuesday-night journal excerpt from Tom v4.3 Rebuild phase)
- *Operating moment events* – Onboarding completed, IDQ intake completed, IDQ retake completed, False Start registered, Cycle 2 onset reached

Loading and retrieval

Loaded per draft request. Same view-based access pattern as Layer 3. The view exposes only the last 90 days by default; longer windows are accessible via parameterized retrieval if the operating moment template requires it (e.g., a Cycle 2 Welcome draft may pull Cycle 1's full history).

Update and persistence rules

Read-only. The Founder Agent does not write to Layer 4. Member Agent writes propagate immediately.

Voice and governance enforcement at this layer

Layer 4 is what makes the drafts feel earned rather than generic. When the Retake Commentary draft references the cost-is-me line from Reshma's Day 50 Disinformation Audit closing reflection, the line is in the Layer 4 context for that draft request. The agent did not invent the line; the agent quoted what the Member already wrote. Voice fidelity at the moment of reflection requires Member-history fidelity.

Failure modes

Excerpt cherry-picking – the agent retrieves a Layer 4 excerpt that does not match the operating moment's tone (e.g., a celebratory Pulse Check note in a False Start Return draft). Mitigation: per-operating-moment retrieval scoping; Jay's review catches misfires before send.

Privacy boundary – the agent surfaces a Conversation Surface excerpt the Member would not expect to see quoted in an email. Mitigation: the Member Agent flags excerpts as Founder-Agent-eligible only when the Member has not marked the session private; Jay's review further filters.

Dependencies

- Member Agent Tech Spec v1.1 §3.4 – the source of Layer 4 in the Member Agent
- Supabase view definition – member_history_for_founder_agent
-

3.5 Layer 5 – Draft Output

Purpose

The per-Member, per-moment artifact the agent produces. The agent's only output type. Lives in the founder_agent_drafts table. Lifecycle: drafted by agent → read by Jay in Review UI → edited by Jay if needed → approved by Jay → sent to Member via HubSpot.

Contents and schema

The founder_agent_drafts table schema:

- `id` (uuid) — primary key
- `member_id` (uuid) — foreign key to members table
- `operating_moment` (enum) — `post_idq_welcome`, `retake_commentary`, `milestone_commentary`, `false_start_return`, `cycle2_welcome`
- `triggered_at` (timestampz) — when the operating-moment trigger fired
- `drafted_at` (timestampz) — when the agent finished generating the draft
- `input_snapshot` (jsonb) — the Layer 3 and Layer 4 context as it existed at draft time (immutable record)
- `corpus_refs` (jsonb) — the Layer 2 retrieval IDs and metadata used to ground the draft
- `draft_subject` (text) — the email subject line
- `draft_body` (text) — the email body, as drafted
- `jay_edits` (text, nullable) — the email body as Jay revised it, if Jay revised
- `approval_status` (enum) — `pending`, `approved`, `rejected`
- `approved_at` (timestampz, nullable) — when Jay approved
- `sent_at` (timestampz, nullable) — when HubSpot confirmed send
- `hubspot_email_id` (text, nullable) — the HubSpot transactional email ID for cross-referencing
- `prompt_version` (text) — the Layer 1 system prompt version used (for voice-regression tracing)
- `library_version` (text) — the Layer 2 corpus version used (for voice-regression tracing)

Storage and persistence pattern

Append-only writes by the agent; updates by Jay through the Review UI. The agent never modifies a draft after writing it; Jay's edits are stored in a separate field. The historical record of every draft, including Jay's edits, persists indefinitely.

HubSpot synchronization

On send, the Review UI calls HubSpot's transactional email API. HubSpot returns an email ID. The Review UI writes the email ID and the sent timestamp to the draft record. The Member's HubSpot contact record now carries the email under their interaction history.

Hand-off mechanisms

On send, the Review UI emits a `Founder-Agent-send-event` to the Member Agent. The Member Agent records the event in its Layer 5 (Conversation Context) and Layer 4 (Member History). The Member Agent's next reflection can reference what Jay just sent. Example: the next IDQ retake reflection might open with "you received Jay's milestone email Tuesday" as continuity.

Voice and governance enforcement at this layer

Layer 5 is where the Jay-review-before-send governance is operationally enforced. The agent writes pending drafts. Drafts do not send. Sending requires Jay's authenticated approval through the Review UI. The architecture is intentionally a one-way gate: drafts flow toward Jay; sends flow away from Jay; the agent does not have the second arrow.

Failure modes

Stuck drafts — a draft sits in pending status for longer than the Jay-review SLA. Mitigation: SLA monitor; Jay-visible queue depth indicator in Review UI; weekly draft-queue review.

HubSpot send failure — the send call fails after Jay approves. Mitigation: send-event retry queue; draft remains in approved-not-sent status until send succeeds; Jay-visible alert if retries exhaust.

Edit-and-resend race — Jay edits a draft after it has already sent. Mitigation: Review UI disables edits on sent-status drafts; new sends require a new draft request.

Dependencies

- Supabase — founder_agent_drafts table
- HubSpot — transactional email API
- Member Agent Tech Spec v1.1 — send-event consumer in Member Agent Layer 4

4. The Six Operating Moments

Each operating moment is a distinct trigger condition, input pattern, posture, output structure, and hand-off path. Five are event-driven and produce per-Member drafts. One is ambient and surfaces library content without per-Member drafting.

4.1 Post-IDQ Welcome

Trigger

Fires when a Member submits their IDQ intake — first IDQ at Cycle 1 entry, or earned/offered IDQ at Cycle 2 onset. Detected via the Member Agent emitting an idq_intake_completed event to the Founder Agent trigger queue.

Inputs

Layer 2 retrieval scoped to: welcome emails Jay has previously written, second-edition Part Three opening register, BP front-matter “welcome to this” constructs. Layer 3: Member ID, name, Door, noun, track, intake quote, cycle. Layer 4: the IDQ intake record itself (numbers, dimension scores, signed deltas if Cycle 2).

Posture

Acknowledge the intake quote in plain language. Name the Door. Frame the cycle ahead. Do not reference metrics or numbers — the welcome is not a results email. The voice is steady and personal, not promotional.

Outputs

A draft email with subject line and body. Subject line pattern: "FirstName — welcome to G4L" for Cycle 1; "FirstName — Cycle 2 begins" for Cycle 2. Body length: 6-10 sentences. The Cycle 1 welcome closes with a single line about the next operating moment (first gating asset). The Cycle 2 welcome closes with a single line about Cycle 2 length and the new Door's lens.

Hand-offs at draft completion

Draft lands in Review UI. Jay-review SLA: 24 hours. On send: HubSpot delivers; Member Agent records the welcome-sent event in Layer 4 for next-reflection context.

Journey dependency

Tom v4.3 §Welcome email; Reshma v1.1 §Welcome email; Martin v1.1 §Welcome email; the second edition Part Three opening passages.

4.2 Retake Commentary

Trigger

Fires when a Member completes an IDQ retake. Detected via the Member Agent emitting an `idq_retake_completed` event with the retake number and the retake's ID Score record.

Inputs

Layer 2 retrieval scoped to: prior retake commentary emails Jay has written, the Founder Agent voice register sections of the canonical journeys at parallel retake numbers. Layer 3: full Member Profile. Layer 4: the new retake record, the previous retake for delta, recent Pulse Check signals (last 60 days), any asset closures since the previous retake (especially the closing reflection text).

Posture

Name what the chart shows in Member-specific terms. Reference the Door explicitly. Quote the Member's own words back when an asset closing reflection or Pulse Check note is striking. Do not assign praise; reflect movement. The dashboard already shows the numbers; the email's job is to name what the numbers mean for this Member.

Outputs

A draft email with subject line and body. Subject line pattern: "FirstName – your Day {N} retake." Body length: 8-14 sentences. The body opens by naming a specific dimension movement, references the Door, optionally quotes the Member's own words, and closes with what is coming next in the cycle. Drafts may include an optional milestone callout if a Phase boundary has been crossed by this retake.

Hand-offs at draft completion

Draft lands in Review UI. Jay-review SLA: 48 hours after the retake – the email should land in the Member's inbox within 2 days of the retake so the retake feels continuous. On send: HubSpot delivers; Member Agent records the retake-email-sent event.

Journey dependency

Tom v4.3 Day 60, 120, 240, 360 retake emails; Reshma v1.1 Day 60 and Day 90 retake emails; Martin v1.1 Day 60, 120, 180, 240, 270 retake emails.

4.3 Milestone Commentary

Trigger

Fires when a gating asset closes that marks a Phase boundary (e.g., Identity Excavation closing in Foundation Phase) OR when a Member-named milestone surfaces (e.g., a Member-reported lab result in a Body-Door journey, a Member-marked life event). Two distinct sub-triggers, same operating moment.

Inputs

Layer 2 retrieval scoped to: milestone emails Jay has written, the Founder Agent voice register at journey milestone moments (Tom's 100-mile milestone, Martin's blood-work labs commentary, Reshma's Disinformation Audit close). Layer 3: full Member Profile. Layer 4: the milestone signal, the Member's closing reflection if asset-driven, the surrounding Conversation Surface context.

Posture

Acknowledge specifically, not generically. Tie to the Door. Tie to the noun if named. Measured tone: a milestone is not a celebration, it is evidence of work.

Outputs

A draft email with subject line and body. Subject line pattern: "FirstName – {milestone in plain language}." Body length: 5-9 sentences. The body acknowledges the milestone, ties it to the Door or noun, references the Member's own words if present, and does not push toward the next thing.

Hand-offs at draft completion

Draft lands in Review UI. Jay-review SLA: 72 hours – milestones can be slightly less time-sensitive than retake commentary. On send: HubSpot delivers; Member Agent records the milestone-email-sent event; optional Circle prompt to the Member if opt-in.

Journey dependency

Tom v4.3 100-mile milestone, Triple Bypass; Reshma v1.1 Disinformation Audit close, Visualization Workshop close; Martin v1.1 Day 150 labs commentary, Athlete naming.

4.4 False Start Return

Trigger

Fires when a Member returns from a missed cadence window. Detected via the Member Agent emitting a member_cadence_resumed event after a member_cadence_window_missed event. The trigger requires the resumption to land within the recovery window defined by the Member Agent.

Inputs

Layer 2 retrieval scoped to: False Start Return emails Jay has written, the Founder Agent voice register at canonical False Start moments. Layer 3: full Member Profile. Layer 4: the miss reason if known, the last activity timestamp, the last Member reflection or asset closure.

Posture

No re-engagement framing. No recovery narrative. The framework holds the gap as a normal feature of a Member's life, not a problem to address. Reference what was happening before – the work practice still in place, the relationships still in place – not what was missed. Steady welcome back.

Outputs

A draft email with subject line and body. Subject line pattern: "FirstName – a note." Body length: 3-6 sentences. The body is short by design. Optionally, a one-line SMS draft for opt-in Members – the SMS draft is also Jay-review-before-send via the Review UI.

Hand-offs at draft completion

Draft lands in Review UI. Jay-review SLA: 24 hours. On send: HubSpot delivers; Member Agent records the False Start Return email-sent event and resumes normal cadence.

Journey dependency

Martin v1.1 Days 60-75 first False Start (Scott's mother's stroke); Martin v1.1 Days 195-220 second False Start (road schedule). The Tom v4.3 master also references a False Start in Rebuild.

4.5 Cycle 2 Welcome

Trigger

Fires once per Member Cycle 2 onset. Detected via the Member Agent emitting a cycle2_onset_completed event with the Cycle 2 trigger type (earned or offered), the new Door if renamed, and the noun (held or compounded).

Inputs

Layer 2 retrieval scoped to: Cycle 2 welcome emails Jay has written. Layer 3: full Member Profile, with both Cycle 1 closing state and Cycle 2 baseline state. Layer 4: full Cycle 1 arc, Cycle 2 IDQ intake record, any inter-cycle Conversation Surface excerpts.

Posture

Name what changed between Cycle 1 close and Cycle 2 onset. Name the new Door. Name the noun (compound if compounded). Frame Cycle 2 length and pacing. Honor the choice to return – the Member did not have to come back.

Outputs

A draft email with subject line and body. Subject line pattern: "FirstName – Cycle 2 begins." Body length: 7-12 sentences. The body opens by acknowledging the Cycle 1 close, names the new Door, names the noun, and closes with the Cycle 2 length framing.

Hand-offs at draft completion

Draft lands in Review UI. Jay-review SLA: 24 hours. On send: HubSpot delivers; Member Agent records the Cycle 2 welcome-sent event; the Member's Cycle 2 cadence begins.

Journey dependency

Tom v4.3 Cycle 2 welcome; Reshma v1.1 §Cycle 2 welcome; Martin v1.1 Cycle 2 welcome.

4.6 Ambient Educational Surfacing

Trigger

Always available. Fires on Member action: dashboard dimension drill-down, asset view with Founder framing slot, or any other dashboard or asset surface that requests Founder-voice content.

Inputs

Layer 3 Member Profile: current Door, current dimension state, current Phase. The agent queries the founder_agent_library table for pre-approved Founder-voice cards matching (Door, dimension, Phase) and surfaces the best-match card. No live generation.

Posture

Library-only. The card is pre-approved by Jay during library curation. The agent does not draft live; the agent retrieves and surfaces.

Outputs

A Founder-voice card. Length: typically 80-160 words. Rendered on the dashboard's educational layer in the Founder-voice register defined by Brand Standards v1.3.

Hand-offs at surfacing

No hand-off to Jay; the card was pre-approved at library write time. The dashboard records the surfacing event for analytics.

Library curation cadence

The founder_agent_library is curated quarterly by Jay. Each card carries (Door, dimension, Phase) targeting, a body field, and an approved_at timestamp. New Door categories (e.g., when the canonical persona set expands) require library expansion. Cards can be deactivated but not deleted; the historical record persists.

Journey dependency

Tom v4.3 §Door integration card text; Reshma v1.1 §Door integration card text; Martin v1.1 §Door integration card text. These card excerpts are the canonical library seed.

5. Surfaces, Voice, and Governance

5.1 Email as primary surface

The Founder Agent's primary Member-facing surface is email. The five event-driven operating moments all produce email drafts. Email is the rail Jay personally controls — he reads each draft, edits where needed, and clicks send. The Member experiences the email as Jay writing to them, because functionally that is what is happening: Jay read and approved the prose before it left.

Email surface affordances. Subject line: Member's first name + plain-language operating-moment cue ("welcome to G4L," "your Day 60 retake," "a note"). Sender: Jay Crain. Reply-to: Jay's personal G4L address (replies surface back to Jay, not to the Founder Agent). Body: plain text, no marketing markup, no tracking pixels by default.

5.2 Dashboard educational layer as secondary surface

The dashboard's educational layer is the Founder Agent's secondary surface. When a Member drills into a dimension (e.g., taps the Self bar on the four-bar chart) or opens an asset that has a Founder-framing slot, the dashboard surfaces a Founder-voice card. The card is pre-approved library content, not a live draft.

Educational layer affordances. Card render: bordered region in the Founder-voice color register defined by Brand Standards v1.3. Card body: 80-160 words of Founder-voice prose targeting (Door, dimension, Phase). No interactive elements within the card. No CTA. The card is read-only narrative scaffolding for the Member's drill-down.

5.3 Voice register — Jay's voice

The Founder Agent's voice is Jay's voice. The voice is steady, specific, non-promotional, non-hyperbolic. The agent does not write in the second-person motivational register that pervades the wellness category. Per Brand Standards v1.3 voice rules: no exclamation marks except in rare structural emphasis (typically once per email, maximum); no italicized words for emphasis; no questions used as rhetorical devices; no "you've got this" constructions; no body-shaming language even in inverse; no medical claims; no time-pressured calls to action.

Voice fingerprints. Sentence rhythm: short-medium-short. Paragraph length: 2-4 sentences. Signature lines: "He signs with his first name" — the canonical journey-master idiom that describes Jay's actual signoff. Reflective register: name a thing specifically; name what it means in this Member's case; do not generalize across the program.

Voice retrieval grounding. Layer 2 retrievals are the substrate that makes the voice durable across thousands of drafts. The agent's drafts read like Jay because Jay's prior writing is in the inference context at draft time — not because the agent has "learned to write like Jay" in an abstract sense.

5.4 Jay-review-before-send mandatory governance

The canonical governance posture: every Member-facing email Jay sends is one Jay personally read, approved, and sent. The Founder Agent does not autonomously send. The architecture makes this physically impossible – the agent has no send tool, no HubSpot-write tool, no email-API tool. The only write path the agent has is to the `founder_agent_drafts` table in pending status. The send path is the Review UI, gated by Jay's authentication.

Operating implications. The Founder Agent's throughput is bounded by Jay's review time, not by inference speed. At Member count one thousand, Jay's review workload is approximately four to six thousand drafts annually – roughly fifteen to twenty drafts per business day at steady state. The Review UI is designed for this throughput: drafts queued by operating-moment and Member, side-by-side Member context, in-place editing, single-click send. Jay's review time per draft averages 90 seconds to 4 minutes depending on operating moment and Member complexity.

What Jay-review-before-send is not. It is not a rubber stamp. Jay is expected to edit drafts that miss tone, miss Member context, or miss the moment. Jay's edits become Layer 5 (Draft Output) record. Edit patterns inform corpus curation. Drafts that consistently require heavy editing flag voice drift or input-data issues.

5.5 What the Founder Agent does not do

Explicit prohibitions on the Founder Agent, beyond the no-autonomous-send mandate:

- No conversation surface. The Founder Agent does not chat with Members. Members who want conversational interaction interact with the Member Agent. Members who reply to a Jay-sent email reach Jay's inbox, not the Founder Agent.
- No coaching or clinical advice. Drafts may reflect, name, acknowledge, and frame, but do not prescribe behavior change protocols, diet plans, exercise plans, or therapeutic advice. The Greg sub-corpora are scoped to the Member Agent's educational layer for this reason.
- No fabrication of Member data. Drafts reference only Layer 3 and Layer 4 facts. If a draft requires reference to data not in the Member's record, the agent must omit the reference – it cannot infer or invent.
- No autonomous corpus updates. Layer 2 (Jay's writing corpus) grows only by Jay's explicit signoff. The agent's own drafts – even ones Jay approves and sends – do not auto-join Layer 2.
- No outbound reach beyond drafted-and-Jay-sent emails. The agent does not initiate Member contact. The agent does not push notifications, schedule SMS sends, post in Circle, or take any other Member-visible action. Drafts to Review UI is the agent's entire action space.

6. Build Sequence – Single October Launch

Per the May 2026 cascade, the platform launches in a single October 2026 release. There is no June or September staged activity. The Founder Agent's build sequence aligns with the Member Agent's and the broader platform build.

6.1 Pre-October – prep window

From mid-June 2026 through end of September 2026: roughly fifteen weeks of pre-launch build for the Founder Agent.

- *Weeks 1-2* (mid-June to end June): Layer 1 system prompt drafted and Jay-reviewed. Layer 2 corpus ingestion pipeline scaffolded. Brand Standards v1.3 voice rules locked.
- *Weeks 3-5* (July): Layer 2 corpus initial ingest – Jay's prior emails, second edition Part Three, BP v6.1 front matter, three canonical journey Founder-voice excerpts. Vector embedding pipeline live. `founder_agent_corpus` table populated.
- *Weeks 6-8* (August): Layer 3 and Layer 4 Supabase views defined. Member Agent emits events to a development trigger queue. Founder Agent draft pipeline produces drafts against development data.
- *Weeks 9-11* (September): Review UI built and Jay-tested. Founder Agent draft pipeline produces drafts against staging-data Members. Jay reviews staging drafts end-to-end. Voice quality calibration.
- *Weeks 12-15* (late September through pre-launch): Library curation – initial `founder_agent_library` populated with Door + dimension + Phase cards seeded from journey master excerpts. HubSpot transactional email integration verified. Send pipeline tested end-to-end.

6.2 October launch – surface set on Day 1

At launch, the Founder Agent operates the following surfaces:

- *Post-IDQ Welcome* – fires for every Member completing IDQ intake. Day 1 of every Member's Cycle 1.
- *Ambient Educational Surfacing* – dashboard educational layer cards available immediately. Initial library covers all three canonical journey Doors (Career, Aging Parents, Body) across all four dimensions and three Phases.

The other event-driven moments – Retake Commentary, Milestone Commentary, False Start Return, Cycle 2 Welcome – do not fire in the first 30 days of any Member's cycle. They become operationally live as Members reach the corresponding milestones.

6.3 Cycle 1 maturation – first 90 days post-launch

- *Day 30: First Pulse Checks* fire from the Member Agent. The Founder Agent's Retake Commentary moment does not yet fire – retakes are Day 60 minimum.
- *Day 60: First Retake Commentary* drafts begin landing in Jay's Review UI. Retake Commentary is the highest-volume operating moment over the year; expect 1-2 drafts per day at steady state.
- *Day 90: Reshma-track Members* reach closing retake. First Cycle 2 Welcome drafts begin landing for 3-month-track Members.
- *Throughout:* Milestone Commentary and False Start Return drafts land event-driven, irregularly.

6.4 Voice quality monitoring – ongoing

Every draft Jay approves carries a voice-quality flag (none, minor edit, major edit) recorded at send time. The flags feed a weekly voice-quality dashboard Jay reviews. Patterns of major edits on a specific operating moment, Door, or Member archetype indicate voice drift or input-data issues. Voice-quality monitoring is a Jay-owned operational practice, not an automated remediation – the corrective action is system prompt tuning or corpus expansion, both of which require Jay's explicit signoff.

7. Schema Requirements

7.1 founder_agent_drafts table

Full schema documented in §3.5. Summary: one row per draft. Fields cover Member ID, operating moment, trigger time, draft time, input snapshot, corpus refs, draft subject and body, Jay's edits, approval status, approval and send timestamps, HubSpot email ID, system prompt version, corpus version.

Indexes: member_id + operating_moment + triggered_at for queue filtering; approval_status for Review UI queue depth; sent_at for send-time analytics.

7.2 founder_agent_corpus table

Layer 2 storage. One row per Jay-signed-off artifact. Fields: id (uuid), artifact_text (text), source (enum: email, book_part_three, bp_front_matter, brand_standards, circle_post, journey_excerpt), source_metadata (jsonb), original_date (date), embedding (vector), signoff_status (enum: pending, approved, deprecated), signoff_at (timestampz), curator (text, default 'Jay Crain'), corpus_version (text).

Indexes: signoff_status for retrieval filtering; embedding for vector similarity search; source for source-faceted retrieval.

7.3 founder_agent_library table

Ambient layer storage. One row per pre-approved Founder-voice card. Fields: id (uuid), door (text), dimension (text: physical, self, social, outlook), phase (text: foundation, rebuild, sustain), card_title (text, optional), card_body (text), approved_at (timestampz), approved_by (text, default 'Jay Crain'), active (boolean), surfacing_count (int, incremented on serve), library_version (text).

Indexes: door + dimension + phase + active for surfacing query; approved_at for curation cadence tracking.

7.4 Cross-table relationships

- founder_agent_drafts.member_id → members.id (Member Agent canonical)
- founder_agent_drafts.corpus_refs → founder_agent_corpus.id (array, via input_snapshot.corpus_refs[])
- founder_agent_drafts.hubspot_email_id → HubSpot transactional email record (external system)
- founder_agent_library reads no foreign keys; cards target (door, dimension, phase) tuples, not specific Members

7.5 Read views into the Member Agent's data

Two Supabase views own the cross-agent data contract:

- member_profile_for_founder_agent — read-only projection of the Member Agent's Layer 3, scoped to Member identity, opt-ins, Door, noun, track length, intake quote, cycle, Phase, and Cycle 2 trigger type.
- member_history_for_founder_agent — read-only projection of the Member Agent's Layer 4, scoped to IDQ retake records, Pulse Check signals, asset closures with closing-reflection text, Conversation Surface excerpts the Member Agent flagged as Founder-eligible, and operating moment events.

Views are owned by the Member Agent team; updates to view definitions require coordinated Member Agent + Founder Agent version bumps. Default view window is 90 days; longer windows accessible via parameterized retrieval at Cycle 2 transitions.

8. AI Governance and Prohibitions

The Founder Agent operates under the AI Governance Framework v1.2, with additional Founder-specific provisions documented in this section.

8.1 No autonomous send – canonical

Repeated for emphasis: the Founder Agent does not autonomously send any Member-facing communication. The architecture's send path runs through the Review UI gated by Jay's authentication. The agent has no send tool, no email-API tool, no HubSpot-write tool. Sending is not a permission the agent could be granted; it is a capability the agent does not architecturally possess.

8.2 No fabrication of Member data

Drafts reference only data present in Layer 3 (Member Profile Mirror) and Layer 4 (Member History Mirror). The agent does not infer, extrapolate, or invent Member facts. If a draft would require referencing data not in the Mirror, the agent must omit the reference. Implementation: the system prompt names this as an immutable rule; eval suite includes fabrication-detection tests.

8.3 No coaching or clinical advice

Drafts may name, acknowledge, reflect, and frame. Drafts may not prescribe behavior protocols, diet plans, exercise plans, dosage advice, or any other intervention claiming clinical authority. The Greg sub-corpora live in the Member Agent's Layer 2 for surfacing in the educational layer's Science Check Zones – they are not available to the Founder Agent's drafting pipeline. Clinical content's path to a Member is through the Member Agent's asset surface, not the Founder Agent's email surface.

8.4 Crisis disclosure routing

The Founder Agent does not directly handle crisis disclosures – it has no Conversation Surface. If Layer 4 contains a Conversation Surface excerpt flagged by the Member Agent as crisis-relevant (e.g., suicidal ideation language the Member Agent's safety system flagged), the excerpt is masked from the Founder Agent's retrieval. The Founder Agent does not see crisis-flagged content and therefore cannot draft against it. Crisis routing is the Member Agent's responsibility – 988 Suicide and Crisis Lifeline, Direct tier coach notification (if applicable). The Founder Agent's correct behavior in a crisis is non-action; Jay is alerted separately by the Member Agent's safety system.

8.5 AI disclosure

The Member-facing email Jay sends is Jay's email. The Founder Agent's role in the draft is not disclosed within the email body – the email is Jay's name, Jay's voice, Jay's signoff. The fact that drafts originate from an AI drafting system is disclosed at the program's onboarding (AI Governance Framework v1.2 §3) and in the public-facing program documentation. The agent is not hidden – the

agent's existence is named in program materials – but the per-email role is not announced in the email itself, consistent with how an executive's assistant who drafts correspondence is not named in the executive's signature.

8.6 Voice authenticity

The Founder Agent's drafts must read as Jay's voice. The Layer 2 corpus grounding is the structural mechanism; Jay-review-before-send is the operational mechanism. The combination prevents the drafted prose from drifting into generic motivational AI voice. Voice quality monitoring (§6.4) is the longitudinal feedback loop.

8.7 Bias and inclusion

The agent's drafts must work for the full range of Members the canonical journey set demonstrates – Tom Miller (white, married father in Boulder), Reshma Patel (Indian-American daughter in Sacramento), Martin Jackson (Black gay man in Chicago) – and the broader range Members will represent at scale. Voice register that flattens cultural specificity or defaults to a single demographic register is a voice-quality failure. Corpus curation includes inclusion review at quarterly cadence.

9. Open Decisions

Architectural decisions deferred to post-v1.0 review, or to a specific milestone, or to the developer hire to surface during build. Each open decision carries its current default, its decision deadline, and the decision owner.

9.1 Editorial review depth – Jay full vs delegated

Default: Jay personally reviews every event-driven draft at all Member counts in Year 1. Open question: at what Member count does the Jay-review-only model break, and what is the next step – delegated review by a vetted editor working in Jay's voice, or AI-assisted pre-flagging to focus Jay's attention. Deadline: end of Year 1 (Member count target 1000-2000). Owner: Jay.

9.2 Library curation cadence

Default: quarterly Jay review of the founder_agent_library. Open question: is quarterly cadence right, or should library curation be event-driven (e.g., when a new Door category enters the canonical set, when voice-quality monitoring flags drift). Deadline: end of Q1 post-launch. Owner: Jay.

9.3 Ambient surfacing logic

Default: surface the highest-similarity library card for the Member's (Door, dimension, Phase) tuple. Open question: should the surfacing be deterministic (same Member, same drill-down, same card every time) or stochastic (rotate cards within the high-similarity set for variety). Deadline: pre-launch UI build (September 2026). Owner: developer hire to propose; Jay decides.

9.4 SMS draft surface for opt-in Members

Default: SMS drafts on False Start Return are Jay-review-before-send in the same Review UI. Open question: do we extend SMS drafting to other operating moments (Post-IDQ Welcome with a same-day SMS, milestone callouts with optional SMS), and if so, what is the SMS volume budget. Deadline: pre-launch UI build. Owner: developer hire to scope; Jay decides.

9.5 Cycle 2 onset framework distinction – offered vs earned

The current journey set introduces an offered-versus-earned distinction at Cycle 2 onset (Tom and Reshma earn theirs via crisis; Martin's is offered after a stable Cycle 1 close). This distinction is not yet in the BP or the Member Agent Tech Spec v1.1. Open question: does the distinction make it into the canonical framework, and if so, what are the trigger semantics in the Member Agent's Layer 4. Deadline: cascade into Member Agent Tech Spec v1.2 if confirmed. Owner: Jay; team review.

9.6 Persona-spec IDQ baseline interpretation

The personas PDF lists Reshma's IDQ baseline as 43 and Martin's as 14. The journey masters and this Spec read both as 0-120 raw values (so 43 raw / 36 normalized and 14 raw / 12 normalized). If the team intended these as normalized values rather than raw, the raw baselines need correction (approximately 52 and 17 respectively) and the journey trajectories shift. Deadline: this week. Owner: Jay; team confirmation.

9.7 Counsel vs Companion rename

The Member Agent's product naming is a separate open decision – Counsel or Companion or another candidate – carried in the Member Agent Tech Spec v1.1 open decisions. Founder Agent naming is similarly open: Founder Agent is internal name; Member-facing reference is currently “Jay's drafting” in the AI disclosure language. Deadline: pre-launch brand finalization. Owner: Jay.

10. Cross-References

10.1 Documents this Spec depends on

- Member Agent Technical Specification v1.1 – canonical source for Layers 3, 4, and the operating moments the Member Agent emits as triggers
- Brand Standards v1.3 Inputs – voice rules, color register, Founder-voice card visual treatment
- AI Governance Framework v1.2 – prohibitions, AI disclosure, crisis routing
- Tom Miller Member Journey v4.3 master – canonical Founder Agent voice register and operating-moment behavior at all Tom-arc inflection points
- Reshma Patel Member Journey v1.1 master – canonical Founder Agent voice register at 3-month track, Aging Parents Door
- Martin Jackson Member Journey v1.1 master – canonical Founder Agent voice register at 9-month track, Body Door, with the offered Cycle 2 framing
- Business Plan v6.1 – founder voice in BP front matter; second edition Part Three referenced
- Corpus Crosswalk v1 – May 2026 cascade reference
- Platform Development Plan v4 – cross-agent build sequence

10.2 Documents this Spec is referenced by

- Founder Agent Operating Moments SVG v1 – visual companion; §4 in this Spec is the canonical text for the SVG's panel content
- Founder Agent Data Model SVG – visual companion; this Spec's §3 is the canonical text for the data model
- Member Agent Technical Specification v1.1 – Member Agent emits events that this Spec consumes; the contract is bidirectional

10.3 Future documents this Spec anticipates

- Founder Agent Review UI Spec v1 – the Jay-facing draft review surface, not yet documented
- Founder Agent Voice Quality Monitoring Spec v1 – the operational dashboard for ongoing voice quality, scoped in §6.4 but not yet built
- Member Agent Tech Spec v1.2 – carries the offered-vs-earned Cycle 2 distinction if confirmed

11. Document Control

Title: G4L Founder Agent – Technical Specification v1.0

Version: 1.0

Status: Active. v1.0 produced May 2026, cascade applied. Working specification for developer hire, agency RFP, and internal team alignment.

Owner: Jay Crain

Companion documents: Founder Agent Operating Moments SVG v1; Founder Agent Data Model SVG; Member Agent Technical Specification v1.1; Brand Standards v1.3 Inputs; AI Governance Framework v1.2; Tom Miller Member Journey v4.3; Reshma Patel Member Journey v1.1; Martin Jackson Member Journey v1.1; Business Plan v6.1; Corpus Crosswalk v1; Platform Development Plan v4.

Versioning convention: Founder Agent Tech Spec versions in lockstep with the canonical Member journey set. v1.0 maps to Tom v4.3 + Reshma v1.1 + Martin v1.1. Next version bump triggered by cascade or new persona addition to the canonical set.

What changes in v1.1 (anticipated)

v1.1 is anticipated to apply: any May 2026 follow-on cascade decisions; the offered-vs-earned Cycle 2 distinction if confirmed into the BP and Member Agent Tech Spec v1.2; the persona-spec IDQ baseline interpretation if confirmed as normalized rather than raw; any developer-hire-surfaced architectural refinements during the pre-launch build window.

Scope locked in v1.0

Five-layer information model; six operating moments (five event-driven plus one ambient); Jay-review-before-send mandatory governance; HubSpot transactional email rail; Anthropic API provider with abstraction layer; single October 2026 launch; voice quality monitoring as Jay-owned operational practice; founder_agent_drafts, founder_agent_corpus, and founder_agent_library tables; member_profile_for_founder_agent and member_history_for_founder_agent views.